

Fed-GEE: Federated Generalized Estimating Equations for Multisite Health Data, with an Application to Veterans' Health

Feier Chang¹ and Soumik Purkayastha^{1,2*}

¹Department of Biostatistics and Health Data Science, School of Public Health,
University of Pittsburgh, Pittsburgh, Pennsylvania, USA

²Center for Healthcare Evaluation, Research, and Promotion (CHERP),
VA Pittsburgh Healthcare System, Pittsburgh, Pennsylvania, USA

Short title: Small-site inference for federated GEE

Abstract

Joint analysis of multisite health data increases sample size, population diversity, and generalizability. Yet patient records often cannot be pooled, outcomes are correlated within sites, and sites are often few. We propose federated generalized estimating equations (Fed-GEE) for population-averaged associations. By exchanging site score vectors and information matrices, Fed-GEE reproduces the pooled estimate exactly. With few sites as the independent clusters, however, the sandwich variance is too small and normal critical values are too narrow. We derive federated versions of the small-sample variance corrections and reference distributions used in pooled analyses. We give theoretical justification and a simulation-based recommendation: the Kauermann–Carroll correction with Bell–McCaffrey degrees of freedom. In simulations with 10 to 50 sites, Fed-GEE matched pooled GEE to numerical precision, and the recommended procedure kept 95% interval coverage near nominal (0.929 to 0.971), compared with as low as 0.863 without correction. Meta-analysis of site-specific fits covered as low as 0.010. We applied Fed-GEE to 29,041 hospitalizations of 18,900 Veterans with alcohol use disorder, emulating a federation of 18 Veterans Health Administration regional networks. Clinician and non-clinician addiction consultation were each associated with

*Corresponding author: Soumik Purkayastha, Department of Biostatistics and Health Data Science, School of Public Health, University of Pittsburgh, Pittsburgh, PA 15261, USA. Email: soumik@pitt.edu

starting alcohol use disorder medication (odds ratios 1.47 and 1.88). These findings match the pooled analysis, yet no patient records need leave their sites. Non-clinician consultation was also five times as common (15.0% versus 2.9%). Where non-clinician staff cost less to hire, expanding this workforce could increase consultation capacity within a fixed budget.

Keywords: sandwich variance estimator; small-sample correction; few clusters; Bell–McCaffrey degrees of freedom; distributed inference; alcohol use disorder; addiction consultation

1 Introduction

Alcohol use disorder (AUD) is more common among US Veterans than in the general adult population. National surveys estimated lifetime prevalences of 42.2% and 29.1%, respectively [1, 2]. Among Veterans, AUD is associated with substantial psychiatric morbidity, suicidal behavior, and higher all-cause mortality [1, 3].

Each hospitalization for AUD offers an opportunity to initiate treatment [4]. Effective medications for AUD (MAUD) are available but remain underused [5]. Hospital addiction consultation services are associated with greater initiation of medications for opioid use disorder [6, 7]. Evidence for AUD is more limited, although recent studies suggest that consultation may also support MAUD initiation [8, 9].

Using national Veterans Health Administration (VHA) data, we examine whether hospital addiction consultation is associated with higher MAUD initiation and whether this association is present for consultation delivered by both clinicians and non-clinicians. In this study, clinicians are physicians and advanced practice clinicians; non-clinicians are social workers, nurses, and peer support staff [9]. The latter can support treatment engagement and coordination, with prescribing remaining the responsibility of an authorized clinician. These associations can inform staffing policy. Evidence confined to clinician consultation would favor prioritizing clinician capacity; evidence for non-clinician consultation would support considering a broader workforce. Where non-clinician staff are more available and less costly to hire, this could expand consultation capacity within a fixed budget.

We analyze 29,041 hospitalizations of 18,900 Veterans with a primary diagnosis of AUD in 2022 or 2023 [9]. The outcome is MAUD initiation during hospitalization or within 7 days of discharge. Hospitalizations are clustered within Veterans and within regions that share health services and practice patterns. Our policy target is the population-averaged association between consultation and MAUD initiation across the VHA. Generalized estimating

equations (GEE) provide a routine way to estimate this marginal association while allowing correlated outcomes [10, 11].

The scientific setting adds a need to limit the movement of sensitive records. Pooling patient-level data creates additional data-transfer and access risks, and AUD diagnoses carry stigma that can itself discourage treatment seeking [12, 13]. These concerns motivate a federated system in which records remain under local control and sites exchange statistical summaries. Federated methods for generalized linear mixed models are available [14], but their conditional regression coefficients do not directly answer our population-averaged question. We therefore require a federated GEE procedure that supports both marginal estimation and inference. Keeping records local reduces the need to share sensitive data; formal protection against disclosure from exchanged summaries requires additional safeguards [15].

The VHA’s regional organization creates a further inferential challenge. Its 18 Veterans Integrated Service Networks (VISNs) are administrative units responsible for regional resources and policy. Shared staffing policies and practice patterns motivate treating each VISN as an independent unit in our analysis. Thus, despite the large number of hospitalizations, inference rests on only 18 clusters. This does not prevent computation of the GEE point estimate, but it can undermine the large-cluster approximations used for inference, whether estimation is centralized or decentralized. With few clusters, the usual sandwich variance tends to be biased downward and normal Wald intervals can undercover [16, 17]. Established remedies include the Kauermann–Carroll (KC), Mancl–DeRouen (MD), and Fay–Graubard (FG) variance corrections and small-sample reference distributions [16–20]. Their conventional formulations use observation-level residuals and cluster leverage matrices that are directly available in a pooled analysis. Applying them in a federation requires expressions that can be evaluated from shared summaries. The reference distribution must likewise be determined from these summaries, with critical values that reflect unequal information across the small number of sites [19, 21].

The contributions of this paper are as follows. We develop Fed-GEE for population-averaged estimation and small-sample inference in federated networks, with implementations both with and without a coordinating center. We express the KC, MD, and FG variance corrections in score space, using site score vectors and information matrices without transferring patient-level records. We also compute coefficient-specific Bell–McCaffrey degrees of freedom from these summaries, enabling federated construction of the t reference distributions used for confidence intervals and hypothesis tests. Both the variance correction and reference distribution are computed in a single post-convergence stage. With a coordinating center, this requires one additional exchange of site summaries; without a center, the summaries are aggregated through neighbor consensus. Under an affine working-score

model with covariance correct up to scale, we show that the KC half-power gives an exactly unbiased variance estimator and is the unique exponent in the proposed family with this property. Extensive simulations examine coverage with 10 to 50 sites and unequal site sizes, and additional simulations assess decentralized fitting with an uncorrected sandwich. The decentralized small-sample correction is developed theoretically and is not evaluated in these simulations. Finally, we apply the procedure to the VHA staffing question. The application uses centrally available records to emulate federation across VISNs and verify the results against a pooled analysis.

Section 2 presents the VHA analysis model, federated GEE, score-space corrections, and reference distributions, followed by algorithms with and without a coordinating center. Section 3 reports the simulation study, and Section 4 presents the VHA application. Section 5 discusses implications and limitations.

2 Methods

2.1 Application model and estimand

For hospitalization t of Veteran j in VISN i , let Y_{ijt} indicate MAUD initiation within the outcome window defined above. We specify the marginal logistic model

$$\text{logit}\{\Pr(Y_{ijt} = 1 \mid C_{ijt}, N_{ijt}, Z_{ijt})\} = \beta_0 + \beta_C C_{ijt} + \beta_N N_{ijt} + Z_{ijt}^\top \boldsymbol{\gamma}, \quad (1)$$

where C_{ijt} and N_{ijt} indicate clinician and non-clinician consultation, respectively, with no consultation as the reference. The nine adjustment covariates in Z_{ijt} account for measured differences in case mix and hospital care. They cover demographics (age at least 65 years and female sex), alcohol severity and co-occurring substance use (no or low AUDIT-C score and opioid use disorder), medical complexity (Care Assessment Need score at or above the 60th percentile, mild or worse frailty, and intensive care unit stay), and hospital processes (psychiatry discharging service and discharge against medical advice) [9]. These groups correspond to those in Figure 1.

We fit the model by GEE with working independence and $p = 12$ coefficients. Hospitalizations are nested within Veterans and Veterans within VISNs; the VISN is both the unit of federation and the independent unit for variance estimation. The target estimands, $\exp(\beta_C)$ and $\exp(\beta_N)$, are adjusted population-averaged odds ratios for each consultation type relative to no consultation. They describe associations, not causal effects. The following formulation extends this model to general outcomes and working covariances.

2.2 Estimating equations and federated fitting

Suppose the data come from K independent sites. Site i holds n_i patient clusters, indexed by $j = 1, \dots, n_i$. Cluster j has a response vector Y_{ij} , a covariate matrix X_{ij} , and marginal mean $\mu_{ij}(\beta) = g^{-1}(X_{ij}\beta)$ for a known link g and a p -vector β . Let $D_{ij} = \partial\mu_{ij}/\partial\beta^\top$, and let V_{ij} be a working covariance matrix [10]. Each site computes a score vector and an information matrix,

$$\mathcal{S}_i(\beta) = \sum_{j=1}^{n_i} D_{ij}^\top V_{ij}^{-1} \{Y_{ij} - \mu_{ij}(\beta)\}, \quad \mathcal{B}_i(\beta) = \sum_{j=1}^{n_i} D_{ij}^\top V_{ij}^{-1} D_{ij}. \quad (2)$$

The GEE estimator $\hat{\beta}$ solves $\mathcal{S}(\beta) = \sum_i \mathcal{S}_i(\beta) = 0$. Because the score and information matrix are additive across sites, the estimating equation can be solved using site summaries. At the current value $\beta^{(r)}$, each site sends $\mathcal{S}_i\{\beta^{(r)}\}$ and $\mathcal{B}_i\{\beta^{(r)}\}$ to a coordinating center. The center adds them and takes a Fisher-scoring step,

$$\beta^{(r+1)} = \beta^{(r)} + \mathcal{B}\{\beta^{(r)}\}^{-1} \mathcal{S}\{\beta^{(r)}\}, \quad \mathcal{B} = \sum_{i=1}^K \mathcal{B}_i. \quad (3)$$

It then returns $\beta^{(r+1)}$ to the sites. We stop when $\|\beta^{(r+1)} - \beta^{(r)}\| < 10^{-8}$, with at most 50 iterations. Each message has $p + p(p+1)/2$ summary statistics, and no patient-level record leaves a site. Algorithm 1 summarizes the fitting routine and the one-shot post-convergence correction developed in Section 2.5.

2.3 Large-sample inference

Inference uses the cluster-robust sandwich variance at the site level,

$$\hat{V}_S = \hat{\mathcal{B}}^{-1} \hat{\mathcal{M}} \hat{\mathcal{B}}^{-1}, \quad \hat{\mathcal{M}} = \sum_{i=1}^K \hat{\mathcal{S}}_i \hat{\mathcal{S}}_i^\top, \quad (4)$$

where $\hat{\mathcal{S}}_i = \mathcal{S}_i(\hat{\beta})$ and $\hat{\mathcal{B}} = \sum_i \mathcal{B}_i(\hat{\beta})$. The site is the independent unit because patients within a site share practice patterns, staffing, and local policy. The site-level sandwich accommodates dependence among patients within a site and remains asymptotically valid under working-covariance misspecification, provided the marginal mean is correctly specified and the sites are independent [10]. Suppose the sites are independent, p is fixed, $K^{-1}\hat{\mathcal{B}} \rightarrow B_0$, and $K^{-1}\hat{\mathcal{M}} \rightarrow M_0$ in probability. Under the regularity conditions of Liang and Zeger [10], $\hat{\beta}$ is consistent and

$$\sqrt{K}(\hat{\beta} - \beta_0) \xrightarrow{d} N(0, B_0^{-1} M_0 B_0^{-1}) \quad \text{as } K \rightarrow \infty.$$

Hence $\widehat{V}_S$ estimates $\text{Var}(\widehat{\beta})$ directly, and Wald intervals use normal critical values. Web Appendix A states the conditions. The result is standard, because $\widehat{\beta}$ is the pooled GEE estimator.

2.4 The small-site problem

The large-sample approximation depends on the number of independent sites, which can be small even in a large health-system cohort. Estimating the regression coefficients from these same sites constrains the fitted scores and can bias the sandwich variance downward.

Collect the fitted scores as the columns of $S = [\widehat{S}_1, \dots, \widehat{S}_K] \in \mathbb{R}^{p \times K}$. Since $S\mathbf{1}_K = 0$ at the GEE solution,

$$\text{rank}(\widehat{\mathcal{M}}) \leq \min(p, K - 1). \quad (5)$$

Thus, the meat is singular whenever $K \leq p$ (Web Appendix E). For $K > p$, the bound does not imply singularity, but the variance estimate still depends on at most $K - 1$ independent score contrasts.

Fitting also reduces the expected score variation under a correctly specified affine working model. The reduction depends on each site’s contribution to the total information, so unequal information across sites can produce unequal attenuation. This provides a working-model explanation for the downward bias of the sandwich variance with few clusters [16, 17]. The score-space corrections below account for this attenuation through site leverage; Section 2.7 states the conditions under which the correction is exact.

Remark 1 (Independent units in the VHA analysis). The VHA cohort contains 29,041 hospitalizations, but the site-level sandwich is constructed from only 18 VISN scores. Additional hospitalizations can increase information for estimation without increasing the number of independent units available for variance estimation. With $p = 12$ coefficients, the rank bound does not force singularity, but neither the large cohort nor $K > p$ ensures that the large-sample reference distribution is adequate.

2.5 Score-space correction and site leverage

Conventional small-sample corrections use observation-level residuals and cluster leverages [16–18]. In a federation, the coordinating center instead has access to $\widehat{S}_i$ and $\mathcal{B}_i(\widehat{\beta})$. We express the corrections through these summaries, allowing variance estimation to be completed in one additional exchange after convergence.

Define the information-standardized score and site leverage

$$z_i = \hat{\mathcal{B}}^{-1/2} \hat{\mathcal{S}}_i, \quad G_i = \hat{\mathcal{B}}^{-1/2} \mathcal{B}_i(\hat{\beta}) \hat{\mathcal{B}}^{-1/2}. \quad (6)$$

Each G_i is symmetric and positive semidefinite, with $\sum_i G_i = \mathbf{I}$ and $\sum_i \text{tr}(G_i) = p$. It measures site i 's contribution to the total information and serves as the site leverage. For linear regression with one observation per site and $V_{ij} = 1$, $\text{tr}(G_i)$ reduces to the usual hat value h_{ii} . If all sites contribute equal information, $G_i = \mathbf{I}/K$; otherwise, leverage varies with site size and covariate distribution. Web Appendix C establishes invariance to the choice of matrix square root.

We require $\mathbf{I} - G_i \succ 0$ for every site, or equivalently $\hat{\mathcal{B}} - \mathcal{B}_i(\hat{\beta}) \succ 0$. The condition fails only if one site alone identifies some direction of β . We report the smallest eigenvalue of $\mathbf{I} - G_i$ across sites as a diagnostic.

2.6 Variance corrections

Kauermann–Carroll (KC) and Mancl–DeRouen (MD). For an exponent $\alpha > 0$, let $A_i(\alpha) = (\mathbf{I} - G_i)^{-\alpha}$ and

$$\hat{V}(\alpha) = \hat{\mathcal{B}}^{-1/2} \left\{ \sum_{i=1}^K A_i(\alpha) z_i z_i^\top A_i(\alpha) \right\} \hat{\mathcal{B}}^{-1/2}. \quad (7)$$

The adjustment inflates each standardized score according to its site leverage. The choice $\alpha = 1/2$ is the score-space analogue of the KC correction [17]. The choice $\alpha = 1$ is the analogue of the MD correction [16]. We write $\hat{V}_{\text{KC}} = \hat{V}(1/2)$ and $\hat{V}_{\text{MD}} = \hat{V}(1)$. Kauermann and Carroll [17] chose the half-power so that the corrected variance is unbiased in the linear model with a correct working covariance. Mancl and DeRouen [16] derived the full power from a first-order bias approximation. Under the working model of Section 2.7, the full power overestimates the variance (Theorem 1). The adjustment acts on the standardized score. Applying $(\mathbf{I} - G_i)^{-\alpha}$ directly to the raw score $\hat{\mathcal{S}}_i$ gives a different estimator in general (Web Appendix B).

Fay–Graubard (FG). Fay and Graubard [18] rescale each cluster's score by a diagonal factor. In score space the factor for site i is

$$F_i = \text{diag} \left[\left\{ 1 - \min_{k=1} \left(b, (\mathcal{B}_i \hat{\mathcal{B}}^{-1})_{kk} \right) \right\}^{-1/2} \right]^p, \quad \hat{V}_{\text{FG}} = \hat{\mathcal{B}}^{-1} \left(\sum_{i=1}^K F_i \hat{\mathcal{S}}_i \hat{\mathcal{S}}_i^\top F_i \right) \hat{\mathcal{B}}^{-1}, \quad (8)$$

with the default bound $b = 0.75$. FG inflates each coordinate of the score separately using the diagonal of the site's information share. The bound limits inflation when information about a coefficient is concentrated in one site. Li and Redden [20] found FG preferable to

201 KC when cluster sizes vary widely.

202 All three estimators share a common form. Write the corrected raw score as $T_i \hat{\mathcal{S}}_i$, with
 203 $T_i = \hat{\mathcal{B}}^{1/2} A_i(\alpha) \hat{\mathcal{B}}^{-1/2}$ for KC and MD and $T_i = F_i$ for FG. Then each estimator equals

$$\hat{V}_T = \hat{\mathcal{B}}^{-1} \left(\sum_{i=1}^K T_i \hat{\mathcal{S}}_i \hat{\mathcal{S}}_i^\top T_i^\top \right) \hat{\mathcal{B}}^{-1}. \quad (9)$$

204 The uncorrected sandwich (4) is the case $T_i = \mathbf{I}$.

205 *Remark 2* (Relation to CR2). With working independence and an identity link, $\hat{V}_{\text{KC}}$ equals
 206 the bias-reduced linearization estimator CR2 [19, 22]. For logistic GEE the two differ,
 207 because the working weights are not constant. In the VHA application they differed by less
 208 than 0.5% (Section 4).

209 2.7 Exactness of the half-power

210 The half-power in the KC correction is exact under a working model for the scores. The
 211 model treats $\hat{\mathcal{B}}$ and the $\mathcal{B}_i$ as fixed matrices B and B_i .

212 **Theorem 1** (Affine working-score exactness). *Conditionally on the design, suppose $\mathcal{S}_i(\beta) =$
 213 $U_i - B_i(\beta - \beta_0)$, where the B_i are fixed positive semidefinite matrices, the U_i are independent,
 214 $E(U_i) = 0$, and $\text{Var}(U_i) = \tau B_i$ for some $\tau > 0$. Suppose $B = \sum_i B_i \succ 0$ and $B - B_i \succ 0$ for
 215 every i . Define G_i and z_i by (6) with $\hat{\mathcal{B}}$ replaced by B . Then*

$$E(z_i z_i^\top) = \tau G_i (\mathbf{I} - G_i), \quad E \left\{ \sum_{i=1}^K A_i(\tfrac{1}{2}) z_i z_i^\top A_i(\tfrac{1}{2}) \right\} = \tau \mathbf{I}.$$

216 Hence $\hat{V}_{\text{KC}}$ is unbiased for $\text{Var}(\hat{\beta}) = \tau B^{-1}$ at every finite K . Within the family $A_i(\alpha)$,
 217 $\alpha = 1/2$ is the only exponent that is unbiased for all admissible $\{B_i\}$. The choice $\alpha = 1$
 218 overestimates the variance.

219 The proof is given in Web Appendix B. Under the affine working model, the half-power
 220 adjustment removes the attenuation in the expected outer product of each standardized
 221 score. The condition $\text{Var}(U_i) = \tau B_i$ requires the working covariance to be correct up to a
 222 common scale, as in the original KC derivation [17]. The result establishes exactness and
 223 uniqueness of the exponent within this family under that condition. For logistic GEE or
 224 a misspecified working covariance, the correction is based on linearization; its finite-sample
 225 performance is evaluated in Section 3.

226 *Remark 3.* As $K \rightarrow \infty$ with no dominant site, $\max_i \lambda_{\max}(G_i) \rightarrow 0$ in probability. Then
 227 $\hat{V}_{\text{KC}}$, $\hat{V}_{\text{MD}}$, $\hat{V}_{\text{FG}}$, and $\hat{V}_{\text{S}}$ have the same first-order limit (Web Appendix D). Their effects are

therefore asymptotically negligible under these conditions.

2.8 Reference distribution

Small-sample inference also requires an appropriate reference distribution. The Wald statistic $(\hat{\beta}_j - \beta_j)/\hat{V}_{jj}^{1/2}$ is compared with a reference distribution, and the normal distribution is too light-tailed when K is small [20]. Common choices are the t distribution with $K - p$ or $K - 1$ degrees of freedom [16, 23]. A fixed count does not respond to unequal information across sites. We therefore use the Satterthwaite approximation of Bell and McCaffrey [19], computed in score space.

The variance estimator is constructed from a small number of site scores and can remain highly variable after bias correction [17]. The Bell–McCaffrey degrees of freedom account for this variability by matching the first two moments of the corrected variance under a working model. Equal information across sites gives $K - 1$ degrees of freedom; concentration of information in fewer sites reduces the effective degrees of freedom and increases the critical value.

Fix a coefficient j and let e_j be the j th unit vector. Define $a_i = T_i^\top \hat{\mathcal{B}}^{-1} e_j$. By (9), the corrected variance of $\hat{\beta}_j$ is

$$\hat{V}_{T,jj} = \sum_{i=1}^K (a_i^\top \hat{\mathcal{S}}_i)^2.$$

Under the working model of Theorem 1 with Gaussian U_i , this is a quadratic form in the U_i . Its first two moments are determined by the $K \times K$ matrix

$$C_{ik} = \mathbf{1}\{i = k\} a_i^\top \mathcal{B}_i a_i - (\mathcal{B}_i a_i)^\top \hat{\mathcal{B}}^{-1} (\mathcal{B}_k a_k), \quad i, k = 1, \dots, K. \quad (10)$$

Matching these moments to a scaled χ^2 distribution gives the degrees of freedom

$$\nu_j^{\text{BM}} = \frac{\{\text{tr}(C)\}^2}{\text{tr}(C^2)}. \quad (11)$$

The interval for β_j is $\hat{\beta}_j \pm t_{\nu_j^{\text{BM}}, 0.975} \hat{V}_{T,jj}^{1/2}$.

The degrees of freedom are computed from the fitted information matrices $\mathcal{B}_i$ and correction operators T_i , without additional outcome summaries. For nonlinear models, these matrices can depend on the outcomes through $\hat{\beta}$. When all sites carry equal information, $\nu_j^{\text{BM}} = K - 1$. When information concentrates in a few sites, ν_j^{BM} falls below $K - 1$, and it can differ across coefficients. This matches the design-based logic of Bell and McCaffrey [19] and Imbens and Kolesár [21]. With working independence and an identity link, ν_j^{BM} equals the CR2 Satterthwaite degrees of freedom of Pustejovsky and Tipton [22]. Web Appendix

F gives the derivation. For FG we also report the degrees of freedom proposed by Fay and Graubard [18] as a sensitivity analysis. Our recommended procedure is $\hat{V}_{KC}$ with a $t_{\nu_j^{BM}}$ reference distribution. Section 3 gives the evidence for this choice.

Remark 4 (Coefficient-specific information across VISNs). As discussed later in the VHA analysis, the Bell–McCaffrey degrees of freedom are 7.4 for non-clinician consultation and 14.6 for psychiatry discharging service (Table 3). Although both coefficients use the same 18 VISNs, information for the consultation contrast is more concentrated across sites. A fixed t_{17} reference would assign the same critical value to both, whereas the coefficient-specific reference assigns a larger critical value to the consultation contrast. The alternative fixed count $K - p = 6$ would give wider intervals than the Bell–McCaffrey reference for both coefficients, holding the variance estimates fixed.

2.9 Federated algorithm with a coordinating center

Algorithm 1 combines iterative estimation with the score-space correction and reference distribution. Once fitting converges, each site evaluates its score and information matrix at $\hat{\beta}$ and sends these summaries to the center. The center then computes the corrected variance and coefficient-specific degrees of freedom. This final stage requires no further parameter updates and no additional type of patient-level summary.

Algorithm 1 Fed-GEE with a coordinating center and one-shot small-sample correction

Require: Local datasets at K sites; common initial value $\beta^{(0)}$; working covariance specification; tolerance $\epsilon = 10^{-8}$; maximum iterations $M = 50$; correction KC, MD, or FG.

Ensure: Estimate $\hat{\beta}$, corrected variance $\hat{V}_T$, and coefficient-specific degrees of freedom ν_j^{BM} .

- 1: **Iterative federated fitting**
 - 2: **for** $r = 0, \dots, M - 1$ **do**
 - 3: Center broadcasts $\beta^{(r)}$ to all sites.
 - 4: Each site i computes and sends $\mathcal{S}_i\{\beta^{(r)}\}$ and $\mathcal{B}_i\{\beta^{(r)}\}$ using (2), in parallel.
 - 5: Center sums the scores and information matrices and updates $\beta^{(r+1)}$ using (3).
 - 6: **if** $\|\beta^{(r+1)} - \beta^{(r)}\| < \epsilon$ **then**
 - 7: Set $\hat{\beta} = \beta^{(r+1)}$ and exit the loop.
 - 8: **end if**
 - 9: **end for**
 - 10: If convergence was not reached, report nonconvergence and stop.
 - 11: **One-shot post-convergence correction**
 - 12: Center broadcasts $\hat{\beta}$; each site recomputes and sends $\hat{\mathcal{S}}_i$ and $\mathcal{B}_i(\hat{\beta})$ once.
 - 13: Center forms $\hat{\mathcal{B}}$, checks the leverage conditions in Section 2.5, and constructs the chosen operators T_i from Section 2.6.
 - 14: Center computes $\hat{V}_T$ using (9) and ν_j^{BM} using (10)–(11).
 - 15: **return** $\hat{\beta}$, $\hat{V}_T$, and $\{\nu_j^{BM}\}_{j=1}^p$; form intervals and tests using the corresponding t reference distributions.
-

2.10 Decentralized estimation and inference

When a coordinating center is unavailable, sites can exchange summaries with their neighbors and approximate the required network totals through consensus. Related work by Gu and Chen [24] develops inference for decentralized M-estimation using stochastic-gradient updates and Polyak–Ruppert averaging. Here, we use consensus on GEE scores and information matrices to approximate the Fisher-scoring update and extend the score-space correction to this architecture.

Let $W = (w_{ih})$ be a nonnegative, symmetric, doubly stochastic mixing matrix on a connected network, with $w_{ih} = 0$ for non-neighboring sites. Write $J_K = K^{-1}\mathbf{1}_K\mathbf{1}_K^\top$ and assume $\rho = \|W - J_K\|_2 < 1$. For a collection of local vectors or matrices $Q_1, \dots, Q_K$, define

$$\mathcal{A}_{L,i}(Q) = \sum_{h=1}^K [W^L]_{ih} Q_h. \quad (12)$$

This quantity is computed through L rounds of neighbor exchange, with $Q_i^{[\ell+1]} = \sum_h w_{ih} Q_h^{[\ell]}$. As L increases, $\mathcal{A}_{L,i}(Q)$ approaches the network average, and $K\mathcal{A}_{L,i}(Q)$ approaches the sum [25]. For symmetric W , $\|W^L - J_K\|_2 = \rho^L$; the corresponding maximum row-sum error is bounded by $\sqrt{K}\rho^L$. Communication requirements therefore depend on both network connectivity and the number of sites.

Each site maintains a parameter vector $\beta_i^{(r)}$. At each iteration, sites first average their parameter states through consensus, evaluate their local scores and information matrices, and then average these summaries. Writing the resulting quantities as $\tilde{\beta}_i^{(r)}$, $\bar{\mathcal{S}}_i^{(r)}$, and $\bar{\mathcal{B}}_i^{(r)}$, the update is

$$\beta_i^{(r+1)} = \tilde{\beta}_i^{(r)} + \eta_r \{\bar{\mathcal{B}}_i^{(r)}\}^{-1} \bar{\mathcal{S}}_i^{(r)}, \quad (13)$$

where $\eta_r > 0$ is a step size. The factor K cancels between the averaged score and information matrix. After terminal parameter consensus, denote the local estimates by $\hat{\beta}_i$ and their network average by $\hat{\beta}_{\text{dec}}$.

Web Appendix H establishes consistency when parameter-consensus error, score-consensus error, and the terminal estimating-equation residual vanish. If their combined error is $o(K^{-1/2})$, then $\hat{\beta}_{\text{dec}}$ has the same first-order distribution as the estimator obtained with a coordinating center. These conditions require sufficient communication and numerical convergence. With $W = J_K$, a single round gives exact averages and, under common initialization and unit step sizes, reproduces the coordinated update.

The small-sample correction is performed after fitting. Each site evaluates $\mathcal{S}_i(\hat{\beta}_i)$ and $\mathcal{B}_i(\hat{\beta}_i)$, then uses consensus to obtain an approximation $\hat{\mathcal{B}}^{(i)}$ to the total information. From this matrix and its local information, site i constructs its correction operator T_i and corrected

meat contribution $T_i \hat{\mathcal{S}}_i \hat{\mathcal{S}}_i^\top T_i^\top$. A further consensus stage aggregates these contributions. Under the smoothness and eigenvalue conditions in Web Appendix H, the corrected variance approaches its coordinated counterpart as parameter and summary consensus errors vanish.

The Bell–McCaffrey degrees of freedom can also be computed from additive summaries. For each coefficient, let $a_i = T_i^\top \hat{\mathcal{B}}^{-1} e_j$, $d_i = a_i^\top \mathcal{B}_i a_i$, and $u_i = \mathcal{B}_i a_i$. The traces in (11) depend only on $\sum_i d_i$, $\sum_i d_i^2$, $\sum_i u_i u_i^\top$, and $\sum_i d_i u_i u_i^\top$, together with $\hat{\mathcal{B}}^{-1}$. Web Appendix H gives the identities. Consensus approximations to these sums therefore provide coefficient-specific reference distributions without collecting patient-level records.

Algorithm 2 summarizes this extension. The correction is one-shot in the sense that it is a single post-fitting stage; it requires sequential consensus operations, each of which may involve several neighbor exchanges.

Algorithm 2 Decentralized Fed-GEE with a post-convergence score-space correction

Require: Local datasets; mixing matrix W ; initial states $\beta_i^{(0)}$; step sizes η_r ; consensus budgets $L_\beta, L_S, L_B, L_M, L_\nu$; stopping tolerances and iteration limit; correction KC, MD, or FG.

Ensure: Local approximations to the coordinated estimate, corrected variance, and Bell–McCaffrey degrees of freedom.

- 1: **Iterative decentralized fitting**
 - 2: **for** $r = 0, 1, \dots$ up to the iteration limit **do**
 - 3: Each site sets $\tilde{\beta}_i^{(r)} = \mathcal{A}_{L_\beta, i}(\beta^{(r)})$.
 - 4: Each site computes $\mathcal{S}_i(\tilde{\beta}_i^{(r)})$ and $\mathcal{B}_i(\tilde{\beta}_i^{(r)})$ using (2).
 - 5: Apply $\mathcal{A}_{L_S, i}$ to scores and $\mathcal{A}_{L_B, i}$ to information matrices to obtain $\bar{\mathcal{S}}_i^{(r)}$ and $\bar{\mathcal{B}}_i^{(r)}$.
 - 6: Update each $\beta_i^{(r+1)}$ using (13).
 - 7: At a candidate stopping point, perform terminal parameter consensus and recompute score and information averages.
 - 8: Stop when all sites satisfy the terminal residual and agreement tolerances of Web Appendix H, using neighbor communication to confirm.
 - 9: **end for**
 - 10: If convergence was not reached, report nonconvergence and stop.
 - 11: **Post-convergence correction**
 - 12: Each site computes $\hat{\mathcal{S}}_i = \mathcal{S}_i(\hat{\beta}_i)$, $\mathcal{B}_i = \mathcal{B}_i(\hat{\beta}_i)$, and $\hat{\mathcal{B}}^{(i)} = K \mathcal{A}_{L_B, i}(\mathcal{B})$.
 - 13: Check leverage conditions and construct T_i using $\hat{\mathcal{B}}^{(i)}$ and $\mathcal{B}_i$ as in Section 2.6.
 - 14: Form $M_i = T_i \hat{\mathcal{S}}_i \hat{\mathcal{S}}_i^\top T_i^\top$ and set $\hat{V}_T^{(i)} = (\hat{\mathcal{B}}^{(i)})^{-1} \{K \mathcal{A}_{L_M, i}(M)\} (\hat{\mathcal{B}}^{(i)})^{-1}$.
 - 15: For each coefficient j , form d_i, u_i using $a_i = T_i^\top (\hat{\mathcal{B}}^{(i)})^{-1} e_j$; apply $K \mathcal{A}_{L_\nu, i}$ to $d_i, d_i^2, u_i u_i^\top, d_i u_i u_i^\top$.
 - 16: Compute local trace approximations and $\nu_j^{\text{BM}, (i)}$ using Web Appendix H; increase consensus budgets if the required positivity conditions fail.
 - 17: **return** $\hat{\beta}_i, \hat{V}_T^{(i)}$, and $\{\nu_j^{\text{BM}, (i)}\}_{j=1}^p$ at each site.
-

3 Simulation study

The simulations compare variance corrections and reference distributions when sites are few. We varied the number of sites, the model dimension, the outcome prevalence, and the balance of site sizes. Unequal site sizes are common in health systems. They lower the Bell–McCaffrey degrees of freedom but leave fixed counts such as $K - 1$ and $K - p$ unchanged. We also compared Fed-GEE with pooled and meta-analytic GEE.

3.1 Design

We generated clustered binary outcomes with a marginal logistic mean. Visit t of patient j at site i had a covariate $x_{ijt} \sim N(0, 1)$ and marginal probability $\text{logit}^{-1}(\beta_0 + \beta_x x_{ijt})$, with $\beta_x = 1$. Outcomes were produced by thresholding a latent Gaussian variable,

$$Z_{ijt} = \sqrt{\rho_H} U_i + \sqrt{\rho_P - \rho_H} W_{ij} + \sqrt{1 - \rho_P} \varepsilon_{ijt},$$

with independent standard normal U_i , W_{ij} , and ε_{ijt} . This construction fixes the marginal mean exactly and induces correlation within patients ($\rho_P = 0.30$) and within sites ($\rho_H = 0.05$) on the latent scale. Each patient had two to six visits.

The design varied four factors. The number of sites was $K \in \{10, 18, 30, 50\}$, which includes the $K = 18$ of the application. The model dimension was $p \in \{2, 5, 10, 20\}$: an intercept, x , and up to 18 nuisance covariates. When present, the first nuisance covariate was a binary site-level variable. The rest were visit-level continuous and binary variables. We kept scenarios with $K - p \geq 5$, which leaves 13 combinations of K and p . The intercept $\beta_0 \in \{-3, -1, 0\}$ gave average outcome prevalences of 7%, 30%, and 50%. The VHA prevalence of 30.8% lies at the middle level. Site sizes had one of three distributions, all with a mean of about 20 patients. In the equal-size arm, sizes were uniform on 10 to 30 patients, a coefficient of variation (CV) of 0.30. In two unequal arms, sizes were gamma distributed with CV 0.78 and 1.00 and a minimum of five patients. Holding the mean fixed separates the effect of imbalance from the effect of total sample size. The full design had $13 \times 3 \times 3 = 117$ scenarios, each with 1,000 replicates.

Every replicate was fitted by federated GEE with working exchangeability. We computed the uncorrected site sandwich (4) and the KC, MD, and FG corrections of Section 2.6. Each variance was paired with four reference distributions: normal, t_{K-1} , t_{K-p} , and $t_{\nu\text{BM}}$. FG was also paired with its own degrees of freedom [18]. We also fitted two comparators. Pooled GEE used all data in one place, with the site-clustered sandwich. Meta-analytic GEE fitted a separate GEE at each site and combined the site estimates of β_x by inverse-variance

weighting. Site-level covariates cannot be estimated within a site and were omitted from the local fits. Sites whose local fit failed or showed separation were excluded.

The target was β_x . We report the coverage of nominal 95% intervals and the ratio of the mean estimated standard error to the empirical standard deviation of $\hat{\beta}_x$ (SE/SD), for which one is ideal. With 1,000 replicates the Monte Carlo standard error of a coverage near 0.95 is 0.007. Table 1 shows the 12 scenarios closest to the VHA application: prevalence 30% and the largest p at each K , which gives the smallest $K - p$ to guarantee non-singularity of the meat matrix. Web Tables S3 and S4 report all 117 scenarios. We summarize results as ranges over scenarios, without averaging.

3.2 Results

The uncorrected sandwich with a normal reference undercovered. Across the 117 scenarios its coverage ranged from 0.863 to 0.952, and 79 scenarios fell below 0.93. The problem was worst with few and unequal sites. At $K = 10$ with CV 0.78 or 1.00, coverage was 0.876 (Table 1). This agrees with earlier work on few clusters [16, 20].

The reference distribution did much of the repair. Keeping the uncorrected variance but using $t_{\nu\text{BM}}$ raised coverage to 0.920 to 0.962. The remaining gap came from the biased variance, with SE/SD as low as 0.81. Adding the KC correction closed it. KC with $t_{\nu\text{BM}}$ covered at 0.929 to 0.971, and only one scenario fell below 0.93. It was within 0.01 of 0.95 in 94 of the 117 scenarios. The corresponding counts were 74 for KC with t_{K-p} , 67 for KC with t_{K-1} , 72 for MD with $t_{\nu\text{BM}}$, and 16 for the uncorrected sandwich with a normal reference. The range was similar at each prevalence: 0.929 to 0.965, 0.935 to 0.971, and 0.938 to 0.969 for prevalences 7%, 30%, and 50%. In the scenario closest to the application ($K = 18$, $p = 10$), it covered at 0.949 to 0.958 across the three site-size distributions, compared with 0.902 to 0.933 for the uncorrected sandwich.

The choice of degrees of freedom mattered consequentially when sites were unequal. With equal sites, ν^{BM} was close to $K - 1$. With unequal sites it fell well below $K - 1$. At $K = 18$ it was 15.3, 10.9, and 9.0 for CV of 0.30, 0.78, and 1.00. A fixed t_{K-1} reference ignores this and undercovered with unequal sites, at 0.934 for $K = 10$ with CV of 0.78 and 0.935 for $K = 18$ with CV 1.00. The count $K - p$ also ignores imbalance, and it depends on p . When $K - p$ was small it overcovered, up to 0.973 in Table 1.

The other corrections behaved as the theory suggests. MD overcorrected, with SE/SD from 0.95 to 1.09, as Theorem 1 predicts for the full power. With $t_{\nu\text{BM}}$ it overcovered, up to 0.983. FG was almost identical to KC. The two coverages differed by at most 0.005 in any scenario, and the FG degrees of freedom were close to ν^{BM} (Web Table S3). In our design

the largest site leverage averaged at most 0.44, so the FG bound rarely applied.

Federated and pooled GEE gave the same estimates and standard errors in every replicate. The largest differences were 3×10^{-15} in $\hat{\beta}_x$ and a relative 10^{-8} in the standard error.

Table 2 compares the recommended Fed-GEE procedure with meta-analytic GEE. Fed-GEE covered at 0.929 to 0.971 across all scenarios, and its bias stayed below 0.3 empirical standard deviations. Meta-analytic GEE covered at 0.010 to 0.912. Its standard errors were too small, and its bias grew with K and p , up to 3.68 empirical standard deviations. Each local fit uses about 20 patients. A logistic fit on so few patients is biased away from zero [26]. In our simulations this bias was larger when the local model had more covariates. Averaging many such fits keeps the bias while the standard error shrinks. Sites excluded for failed or separated local fits add to the problem. The largest values of p are therefore a hard setting for meta-analytic GEE. Fed-GEE solves one pooled estimating equation and does not face this issue.

These results suggest the following recommendation: use the KC correction with a $t_{\nu\text{BM}}$ reference distribution. With equal sites, this is close to KC with t_{K-1} . With unequal sites, the Bell–McCaffrey degrees of freedom guard against the undercoverage of a fixed count. The recommendation held for K from 10 to 50, p up to 20, and outcome prevalence from 7% to 50%.

3.3 Decentralized implementation

A separate simulation evaluated the effect of network communication on decentralized estimation and inference (Web Appendix I). It used 5 to 100 sites, four network structures, and 50 rounds per consensus stage. Intervals used the uncorrected site sandwich with a t_{K-p} reference. These results therefore assess consensus approximation rather than the post-convergence small-sample correction in Algorithm 2.

On complete and regional networks, decentralized coverage closely matched the corresponding coordinated analysis. On a single-hub network with $K = 100$, where $\rho^{50} = 0.61$, coverage ranged from 0.06 to 0.58. Although the slope estimate remained within 0.1 of its true value, its error was substantial relative to the estimated standard error. The results show that a fixed communication budget can be adequate for one network and inadequate for another. Web Appendix I reports the full design and results.

4 VHA application

In the cohort described in the Introduction, MAUD initiation occurred in 8,932 hospitalizations (30.8%). Clinician consultation occurred in 834 hospitalizations (2.9%) and non-clinician consultation in 4,354 (15.0%) [9]. We fitted model (1), emulating federation across VISNs using centrally available records and checking the results against a pooled analysis.

The number of Veterans per VISN had a CV of 0.35, close to the equal-size arm of the simulation (CV 0.30). The largest VISN leverage, the largest diagonal element of $\mathcal{B}_i \hat{\mathcal{B}}^{-1}$, was 0.27. The smallest eigenvalue of $\mathbf{I} - G_i$ across VISNs was 0.725, so every correction was well defined and none required regularization. The FG bound was never reached. The closest simulation scenario has $K = 18$, $p = 10$, prevalence 0.30, and CV 0.30 (Table 1). There, the uncorrected sandwich with a normal reference covered at 0.933 and the recommended procedure at 0.958. With more imbalance (CV 0.78) the corresponding values were 0.921 and 0.957.

Table 3 and Figure 1 report the results. Both consultation types were associated with higher MAUD initiation. Clinician consultation had an odds ratio of 1.47 (95% CI, 1.14 to 1.90). Non-clinician consultation had an odds ratio of 1.88 (95% CI, 1.24 to 2.84). The two intervals overlap, and we did not test whether the two odds ratios differ. The results therefore support both consultation types as options for expanding consultation capacity. They do not show that one type is more effective.

The recommended procedure widened every interval relative to pooled GEE with the uncorrected sandwich. On the log scale the widening ranged from 13% to 42%. Most of it came from the reference distribution. The KC correction raised the standard errors by 4% to 12%, and the $t_{\nu\text{BM}}$ critical values were 9% to 28% larger than 1.96. A small correction is expected here. The VISNs are nearly equal in size, so no VISN holds a large share of the information. No association changed statistical significance. The associations with psychiatry discharging service, opioid use disorder, Care Assessment Need score, and frailty remained clear.

The Bell–McCaffrey degrees of freedom ranged from 5.5 to 14.7 across coefficients (Table 3). The lowest values were for discharge against medical advice and non-clinician consultation, whose intervals widened by 42% and 29%, respectively. Remark 4 relates this variation to the choice of reference distribution.

Two checks support the calculation. First, CR2 with Satterthwaite degrees of freedom, computed on the pooled data with the `clubSandwich` package [22], gave standard errors within 0.5% of the federated KC values and the same degrees of freedom to three decimal places. Second, FG with $t_{\nu\text{BM}}$ reproduced the KC interval limits to within 0.01. The un-

corrected sandwich with $t_{\nu\text{BM}}$, FG with its own degrees of freedom, and CR2 gave the same conclusions for every coefficient (Web Table S5).

Meta-analytic GEE gave different point estimates. The odds ratio for female sex was 1.04 (95% CI, 0.88 to 1.24) compared with 1.22 under GEE. The odds ratio for intensive care unit stay was 0.84 (95% CI, 0.72 to 0.99) compared with 0.88. These differences change the reported significance of both contrasts. Meta-analytic GEE does not target the pooled estimand. It weights VISN-specific estimates, and VISNs whose local fit fails may be excluded. Fed-GEE reproduces the pooled estimate exactly and keeps every VISN in the analysis.

5 Discussion

In a federated analysis of 29,041 hospitalizations in 18 VISNs, both clinician and non-clinician addiction consultation were associated with higher MAUD initiation, after accounting for regional clustering and the small number of VISNs. The odds ratios were 1.47 (95% CI, 1.14 to 1.90) and 1.88 (95% CI, 1.24 to 2.84).

The methodological contribution addresses what remains after federated GEE has reproduced the pooled estimate. With few sites, the sandwich variance is too small and normal critical values are too narrow. The standard corrections cannot be computed by a coordinating center in their usual form. We showed that they can be written in terms of the site scores and information matrices alone. The same holds for the Bell–McCaffrey degrees of freedom. The calculation adds one post-convergence exchange of summaries that the sites already send.

The simulations support one procedure: the KC correction with a $t_{\nu\text{BM}}$ reference distribution. Two findings explain why. First, the reference distribution does much of the work. The uncorrected sandwich with $t_{\nu\text{BM}}$ already came close to nominal coverage, and KC closed the remaining gap. This agrees with Kauermann and Carroll [17], who showed that the variability of the sandwich lowers coverage even when its bias is removed. Second, fixed degrees of freedom fail when sites are unequal. The Bell–McCaffrey value equals $K - 1$ for equal sites and falls as information concentrates, while $K - 1$ and $K - p$ do not change. Li and Redden [20] found FG preferable to KC when cluster sizes vary widely. In our design, FG and KC were nearly identical, including at a site-size CV of 1.00. One possible reason is that both corrections are here defined at the site level from the same information shares, but we did not test this explanation.

The application shows the practical value of the procedure. The correction widened every interval, by 13% to 42% on the log scale, but changed no conclusion. The Bell–McCaffrey degrees of freedom identified the contrasts whose information was concentrated in fewer

VISNs. These were non-clinician consultation and discharge against medical advice. Their intervals widened the most. A fixed reference distribution would have treated all contrasts alike.

For VHA policy, the results support non-clinician consultation as a way to expand consultation capacity. Non-clinician consultation was more common than clinician consultation in the cohort, occurring in 15.0% and 2.9% of hospitalizations [9]. Its estimated association with MAUD initiation was larger than that of clinician consultation, although we did not test the difference. These findings agree with the source cohort analysis, which used hospital fixed effects and patient-clustered standard errors [9]. They are also consistent with evidence on addiction consultation services for opioid use disorder [6, 7] and with the more limited evidence for AUD [8]. They are associations. Veterans who received a consultation may differ from those who did not in ways the model does not capture.

Fed-GEE targets marginal, population-averaged associations. A federated mixed model is preferable when the scientific question concerns conditional, site-specific effects [14]. One-shot distributed estimators reduce communication but approximate the pooled estimate [27]. Fed-GEE reproduces it exactly, at the cost of several rounds of communication. Web Appendix G gives a one-step version for settings with a good common starting value. When no coordinating center is available, the decentralized version of Section 2.10 applies. Its main practical demand is enough consensus rounds for the network at hand.

The correction also changes what must be shared. Estimation needs only the sums $\sum_i \mathcal{S}_i$ and $\sum_i \mathcal{B}_i$, which is compatible with secure aggregation. The correction needs each site’s score and information matrix separately. Any added privacy protection must therefore cover site-level summaries, not only their totals. This affects both the communication protocol and the trade-off between privacy and accuracy [15].

This study has limitations. First, the VHA analysis is observational, and the odds ratios are adjusted associations. Second, the records were available centrally, so federation was emulated. This allowed us to verify the federated results against a pooled analysis. Third, the exactness result holds only under an affine working model in which the working covariance is correct up to scale. For logistic GEE the correction is a linearization-based adjustment, and its performance rests on simulation. Fourth, the simulations used working independence, binary outcomes, and one data-generating mechanism with a fixed site-level correlation. Other outcomes, working correlations, and forms of heterogeneity need study. Fifth, we did not evaluate other small-sample approaches, such as the correction of Morel et al. [28] or the wild cluster bootstrap [29]. Finally, federation keeps patient-level records in place but does not by itself provide differential privacy or cryptographic security [12, 15].

In the VHA analysis, both clinician and non-clinician addiction consultation were asso-

ciated with higher MAUD initiation after accounting for regional clustering and the small number of VISNs. These findings support considering a broader consultation workforce while retaining access to clinicians who can prescribe treatment. Fed-GEE provides a way to address such staffing questions across health systems using site summaries: estimation proceeds through iterative communication, followed by a single post-convergence stage for variance correction and reference-distribution calculation. Although motivated by MAUD initiation, the framework applies more broadly to studies of correlated health outcomes where the target is a population-averaged association and patient-level records remain under local control.

Data availability statement

The patient-level VHA data contain sensitive health information and cannot be made publicly available. Access is subject to VHA data-governance and research approval procedures. The R package `FedGEE` implements federated GEE with the score-space corrections and Bell–McCaffrey degrees of freedom described in this paper. It is available at https://github.com/soumikp/2026_fedGee and can be installed with `remotes::install_github("soumikp/2026_fedGee")`. The same repository contains the code to reproduce the simulation study.

Ethics declaration

The institutional review board of the VA Pittsburgh Healthcare System approved the study and granted a waiver of informed consent [9].

Author contributions

Feier Chang: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Visualization, Writing – original draft, Writing – review and editing. Soumik Purkayastha: Conceptualization, Methodology, Resources, Data curation, Supervision, Project administration, Writing – review and editing.

Acknowledgments

We thank Timothy S. Anderson and Utibe R. Essien for access to the data from the source cohort study [9].

Funding

Data curation for the source cohort was supported by the Dr. Eugene Marsh Center for Health Equity Research and Promotion Pilot Program from the U.S. Department of Veterans Affairs and by a Merit Award (I01RD000382; principal investigators, Drs. Anderson and Essien) from the VA Health Systems Research Division. This research was supported in part by the University of Pittsburgh Center for Research Computing and Data, RRID:SCR_022735, through the resources provided. Specifically, this work used the H2P cluster, which is supported by NSF award number OAC-2117681.

Disclaimer

The content is solely the responsibility of the authors and does not necessarily represent the official views of the U.S. Department of Veterans Affairs, the Veterans Health Administration, or the United States Government.

Conflict of interest

The authors declare no conflicts of interest.

Use of generative artificial intelligence

The authors used Claude (Anthropic), through the Claude Code interface, to assist with manuscript drafting, auditing of the analysis and simulation code, and formatting of the FedGEE R package. All output was reviewed and revised by the authors, who take full responsibility for the accuracy, originality, and integrity of the manuscript and the software.

References

- [1] Brian S. Fuehrlein, Natalie Mota, Albert J. Arias, Louis A. Trevisan, Lorig K. Kachadourian, John H. Krystal, Steven M. Southwick, and Robert H. Pietrzak. The burden of alcohol use disorders in US military veterans: Results from the national health and resilience in veterans study. *Addiction*, 111(10):1786–1794, 2016. doi: 10.1111/add.13423.

- [2] Bridget F. Grant, Ris'e B. Goldstein, Tulshi D. Saha, S. Patricia Chou, Jeesun Jung, Haitao Zhang, Roger P. Pickering, W. June Ruan, Sharon M. Smith, Boji Huang, and Deborah S. Hasin. Epidemiology of DSM-5 alcohol use disorder: Results from the national epidemiologic survey on alcohol and related conditions III. *JAMA Psychiatry*, 72(8):757–766, 2015. doi: 10.1001/jamapsychiatry.2015.0584.
- [3] Lydia A. Chwastiak, Robert A. Rosenheck, Rani Desai, and Lewis E. Kazis. The association of psychiatric illness and all-cause mortality in the national department of veterans affairs health care system. *Psychosomatic Medicine*, 72(8):817–822, 2010. doi: 10.1097/PSY.0b013e3181eb33e9.
- [4] Paul J. Joudrey, Benjamin J. Oldfield, Kimberly A. Yonkers, Patrick G. O'Connor, Gretchen Berland, and E. Jennifer Edelman. Inpatient adoption of medications for alcohol use disorder: A mixed-methods formative evaluation involving key stakeholders. *Drug and Alcohol Dependence*, 213:108090, 2020. doi: 10.1016/j.drugalcdep.2020.108090.
- [5] Barbara J. Mason. Looking back, looking forward: Current medications and innovative potential medications to treat alcohol use disorder. *Alcohol Research: Current Reviews*, 42(1):11, 2022. doi: 10.35946/arcr.v42.1.11.
- [6] Itai Danovitch, Samuel Korouri, Harlene Kaur, Gabrielle Messineo, Teryl Nuckols, Waguih W. Ishak, et al. The addiction consultation service for hospitalized patients with substance use disorder: An integrative review of the evidence. *Journal of Substance Use and Addiction Treatment*, 163:209377, 2024. doi: 10.1016/j.josat.2024.209377.
- [7] Allison J. Ober, Cristina Murray-Krezan, Kimberly Page, Peter D. Friedmann, Jess Anderson, Karen Chan Osilla, et al. Hospital addiction consultation service and opioid use disorder treatment. *JAMA Internal Medicine*, 185:624–633, 2025. doi: 10.1001/jamainternmed.2024.8586.
- [8] Sumeet Singh-Tan, Kristine Torres-Lockhart, Andrea Jakubowski, Tiffany Lu, Joanna Starrels, Patricia De Lima, et al. Addiction consult service and inpatient outcomes among patients with alcohol use disorder. *Journal of General Internal Medicine*, 38:3216–3223, 2023. doi: 10.1007/s11606-023-08202-7.
- [9] Timothy S. Anderson, Soumik Purkayastha, Eden Y. Bernstein, Alyssa Parr, Maria K. Mor, Rachel L. Bachrach, Walid F. Gellad, Leslie R. M. Hausmann, Michael J. Fine, and Utibe R. Essien. Initiation of medications for alcohol use disorder among hospitalized

veterans: A retrospective cohort study. *Annals of Internal Medicine*, 179(6):794–803, 2026. doi: 10.7326/ANNALS-26-00089.

[10] Kung-Yee Liang and Scott L. Zeger. Longitudinal data analysis using generalized linear models. *Biometrika*, 73(1):13–22, 1986. doi: 10.1093/biomet/73.1.13.

[11] Scott L. Zeger, Kung-Yee Liang, and Paul S. Albert. Models for longitudinal data: A generalized estimating equation approach. *Biometrics*, 44(4):1049–1060, 1988. doi: 10.2307/2531734.

[12] Nicola Rieke, Jonny Hancox, Wenqi Li, Fausto Milletari, Holger R. Roth, Shadi Albarqouni, Spyridon Bakas, Mathieu N. Galtier, Bennett A. Landman, Klaus Maier-Hein, et al. The future of digital health with federated learning. *npj Digital Medicine*, 3:119, 2020. doi: 10.1038/s41746-020-00323-1.

[13] K. M. Keyes, M. L. Hatzenbuehler, K. A. McLaughlin, B. Link, M. Olfson, B. F. Grant, and D. Hasin. Stigma and treatment for alcohol disorders in the united states. *American Journal of Epidemiology*, 172(12):1364–1372, 2010. doi: 10.1093/aje/kwq304.

[14] Zhiyu Yan, Kori S. Zachrison, Lee H. Schwamm, Juan J. Estrada, and Rui Duan. A privacy-preserving and computation-efficient federated algorithm for generalized linear mixed models to analyze correlated electronic health records data. *PLOS ONE*, 18(1): e0280192, 2023. doi: 10.1371/journal.pone.0280192.

[15] Peter Kairouz, H. Brendan McMahan, Brendan Avent, Aurélien Bellet, Mehdi Bennis, Arjun Nitin Bhagoji, Kallista A. Bonawitz, Zachary Charles, Graham Cormode, Rachel Cummings, et al. Advances and open problems in federated learning. *Foundations and Trends in Machine Learning*, 14(1–2):1–210, 2021. doi: 10.1561/22000000083.

[16] Lloyd A. Mancl and Timothy A. DeRouen. A covariance estimator for GEE with improved small-sample properties. *Biometrics*, 57(1):126–134, 2001. doi: 10.1111/j.0006-341X.2001.00126.x.

[17] Goran Kauermann and Raymond J. Carroll. A note on the efficiency of sandwich covariance matrix estimation. *Journal of the American Statistical Association*, 96(456): 1387–1396, 2001. doi: 10.1198/016214501753382309.

[18] Michael P. Fay and Barry I. Graubard. Small-sample adjustments for wald-type tests using sandwich estimators. *Biometrics*, 57(4):1198–1206, 2001. doi: 10.1111/j.0006-341X.2001.01198.x.

- [19] Robert M. Bell and Daniel F. McCaffrey. Bias reduction in standard errors for linear regression with multi-stage samples. *Survey Methodology*, 28(2):169–181, 2002.
- [20] Peng Li and David T. Redden. Small sample performance of bias-corrected sandwich estimators for cluster-randomized trials with binary outcomes. *Statistics in Medicine*, 34(2):281–296, 2014. doi: 10.1002/sim.6344.
- [21] Guido W. Imbens and Michal Kolesár. Robust standard errors in small samples: Some practical advice. *Review of Economics and Statistics*, 98(4):701–712, 2016. doi: 10.1162/rest_a_00552.
- [22] James E. Pustejovsky and Elizabeth Tipton. Small-sample methods for cluster-robust variance estimation and hypothesis testing in fixed effects models. *Journal of Business & Economic Statistics*, 36(4):672–683, 2018. doi: 10.1080/07350015.2016.1247004.
- [23] A. Colin Cameron and Douglas L. Miller. A practitioner’s guide to cluster-robust inference. *Journal of Human Resources*, 50(2):317–372, 2015. doi: 10.3368/jhr.50.2.317.
- [24] Jia Gu and Song Xi Chen. Statistical inference for decentralized federated learning. *The Annals of Statistics*, 52(6):2931–2955, 2024. doi: 10.1214/24-AOS2452.
- [25] Lin Xiao and Stephen Boyd. Fast linear iterations for distributed averaging. *Systems & Control Letters*, 53(1):65–78, 2004. doi: 10.1016/j.sysconle.2004.02.022.
- [26] David Firth. Bias reduction of maximum likelihood estimates. *Biometrika*, 80(1):27–38, 1993. doi: 10.1093/biomet/80.1.27.
- [27] Rui Duan, Mary R. Boland, Jason H. Moore, and Yong Chen. ODAL: A one-shot distributed algorithm to perform logistic regressions on electronic health records data from multiple clinical sites. *Pacific Symposium on Biocomputing*, 24:30–41, 2019.
- [28] J.G. Morel, M.C. Bokossa, and N.K. Neerchal. Small sample correction for the variance of GEE estimators. *Biometrical Journal*, 45(4):395–409, 2003. doi: 10.1002/bimj.200390021.
- [29] A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.

Tables

Table 1: Coverage of nominal 95% confidence intervals for β_x in the simulation scenarios closest to the VHA application: outcome prevalence 0.30 and the largest p at each number of sites K . Each row is one scenario with 1,000 replicates. Web Table S3 reports all 117 scenarios.

K	p	CV	Uncorrected		KC			MD	FG	ν^{BM}
			z	$t_{\nu^{\text{BM}}}$	t_{K-1}	t_{K-p}	$t_{\nu^{\text{BM}}}$	$t_{\nu^{\text{BM}}}$	$t_{\nu^{\text{BM}}}$	
10	5	0.30	0.909	0.939	0.950	0.973	0.954	0.966	0.955	8.1
		0.78	0.876	0.933	0.934	0.957	0.951	0.962	0.947	5.8
		1.00	0.876	0.947	0.947	0.965	0.967	0.978	0.967	5.0
18	10	0.30	0.933	0.952	0.956	0.970	0.958	0.965	0.958	15.3
		0.78	0.921	0.950	0.952	0.966	0.957	0.967	0.956	10.9
		1.00	0.902	0.933	0.935	0.952	0.949	0.963	0.946	9.0
30	20	0.30	0.919	0.938	0.942	0.965	0.944	0.949	0.944	26.1
		0.78	0.935	0.956	0.954	0.973	0.964	0.970	0.961	18.2
		1.00	0.927	0.953	0.952	0.964	0.960	0.964	0.960	15.0
50	20	0.30	0.940	0.944	0.946	0.950	0.946	0.950	0.946	44.0
		0.78	0.932	0.943	0.944	0.946	0.949	0.949	0.947	30.5
		1.00	0.934	0.945	0.944	0.946	0.949	0.952	0.949	25.1

CV, coefficient of variation of site size. Uncorrected, site sandwich (4); KC, MD, FG, score-space corrections of Section 2.6. Column headings give the reference distribution. ν^{BM} , median Bell–McCaffrey degrees of freedom for KC. Monte Carlo standard error of each entry is about 0.007.

Table 2: Fed-GEE with the recommended procedure (KC with $t_{\nu_{\text{BM}}}$) and meta-analytic GEE in the scenarios of Table 1. Each row is one scenario. Web Table S4 reports all 117 scenarios.

K	p	CV	Fed-GEE			Meta-analytic GEE		
			Coverage	Bias/SD	SE/SD	Coverage	Bias/SD	SE/SD
10	5	0.30	0.954	0.16	0.988	0.892	0.36	0.858
		0.78	0.951	0.16	0.928	0.870	0.33	0.764
		1.00	0.967	0.21	0.956	0.867	0.36	0.752
18	10	0.30	0.958	0.19	1.012	0.685	1.22	0.844
		0.78	0.957	0.23	1.031	0.757	1.03	0.813
		1.00	0.949	0.19	0.939	0.757	0.86	0.756
30	20	0.30	0.944	0.13	0.978	0.135	2.49	0.741
		0.78	0.964	0.20	0.992	0.297	2.02	0.778
		1.00	0.960	0.23	0.977	0.385	1.81	0.785
50	20	0.30	0.946	0.18	0.984	0.028	3.15	0.724
		0.78	0.949	0.19	0.996	0.110	2.72	0.815
		1.00	0.949	0.15	0.947	0.190	2.32	0.786

CV, coefficient of variation of site size. Bias/SD, bias of $\hat{\beta}_x$ divided by its empirical standard deviation. SE/SD, mean estimated standard error divided by the empirical standard deviation. Meta-analytic GEE uses normal-reference intervals. Pooled GEE is not shown; it reproduced the uncorrected Fed-GEE estimates and standard errors in every replicate.

Table 3: Adjusted odds ratios (OR) and 95% confidence intervals for MAUD initiation in 29,041 hospitalizations of 18,900 Veterans in 18 VISNs.

		Pooled GEE	Fed-GEE (recommended)		Meta-analytic GEE
Covariate	OR	95% CI	95% CI	ν^{BM}	OR (95% CI)
<i>Demographics</i>					
Female	1.22	1.05, 1.41	1.03, 1.44	14.5	1.04 (0.88, 1.24)
Age ≥ 65 years	0.88	0.79, 0.98	0.78, 0.99	13.9	0.86 (0.76, 0.97)
<i>Alcohol severity</i>					
AUDIT-C no or low use	0.89	0.84, 0.93	0.83, 0.94	14.2	0.90 (0.81, 0.99)
Opioid use disorder	0.60	0.54, 0.68	0.53, 0.69	12.4	0.64 (0.57, 0.73)
<i>Medical complexity</i>					
CAN score ≥ 60 th percentile	0.72	0.65, 0.78	0.64, 0.79	14.7	0.73 (0.66, 0.80)
Frailty, mild or worse	0.69	0.61, 0.78	0.60, 0.79	12.7	0.65 (0.55, 0.77)
Any ICU stay	0.88	0.76, 1.02	0.74, 1.05	14.5	0.84 (0.72, 0.99)
<i>Hospital process</i>					
Psychiatry discharging service	1.73	1.43, 2.09	1.39, 2.15	14.6	1.83 (1.68, 1.99)
Addiction consult					
Clinician	1.47	1.19, 1.82	1.14, 1.90	10.4	1.66 (1.45, 1.90)
Non-clinician	1.88	1.36, 2.59	1.24, 2.84	7.4	1.73 (1.22, 2.46)
Discharge against advice	1.13	0.66, 1.92	0.53, 2.41	5.5	1.37 (0.83, 2.27)

Pooled GEE and Fed-GEE give the same OR. Pooled GEE uses the uncorrected VISN-clustered sandwich with a normal reference. Fed-GEE uses the score-space KC correction with a t reference on Bell–McCaffrey degrees of freedom ν^{BM} . Meta-analytic GEE fits a GEE in each VISN and pools the estimates by inverse-variance weighting. CAN, Care Assessment Need; ICU, intensive care unit; AUDIT-C, Alcohol Use Disorders Identification Test–Consumption. Reference categories: male; age 18–64 years; AUDIT-C unhealthy use or higher; no opioid use disorder; CAN score below the 60th percentile; nonfrail or prefrail; no ICU stay; medical or surgical discharging service; no addiction consult; routine discharge.

Figure




Figure 1: Adjusted odds ratios and 95% confidence intervals for MAUD initiation across 29,041 hospitalizations of 18,900 Veterans in 18 VISNs, the units of federation. Fed-GEE uses the score-space Kauermann–Carroll correction with a t reference distribution on Bell–McCaffrey degrees of freedom, computed from the VISN summaries alone. Pooled GEE uses the uncorrected VISN-clustered sandwich with a normal reference. Fed-GEE and pooled GEE give identical odds ratios and differ only in their intervals. Meta-analytic GEE fits a GEE in each VISN and pools the estimates by inverse-variance weighting. Numeric labels give the Fed-GEE odds ratio.